

## **PAGE MODULE:**

- **Page::Page()**  
Zero-initializes an in-memory 8KB page image. Does not set headers yet.
- **void init(page\_no)**  
Formats the page: writes a PageHeader (magic, page\_no) and sets the slotted-page pointers (lower = sizeof(PageHeader), upper = PAGE\_SIZE).
- **void load\_from(raw)**  
Loads a raw 8KB byte buffer (read from disk) into the page's internal data.
- **void store\_to(raw)**  
Writes the current in-memory page image into raw (exactly 8KB) so the caller can persist it.
- **nitems()**  
Returns the current number of **slots** (ItemIds) on the page.
- **free\_space()**  
Returns free bytes available between the ItemId array (lower) and the tuple area (upper). Invariant: free space is [lower, upper).
- **Result add\_tuple(tuple\_bytes)**  
Attempts to insert a variable-length tuple: copies payload at the **top** (decrement upper), appends a new **ItemId** (increment lower).  
Returns the **slot number** on success or NO SPACE if not enough room (needs payload + one ItemId).
- **Result get\_tuple(slot\_no)**  
Fetches the raw tuple payload for a **live** slot (USED and not DEAD). Returns NOT FOUND if slot is out of range or dead/unused.
- **Status mark\_dead(slot\_no)**  
Marks the given slot as **tombstoned** (sets ITEM\_DEAD). Returns OK or NOT FOUND (invalid or unused slot).
- **void compact()**  
Reclaims space by **repacking** all live tuples tightly from the top (reset upper to end of page, rewrite live payloads, rebuild their offsets), and clears dead/unused ItemIds.
- **PageHeader header()**  
Read-only access to the page header (for debugging/inspection: see magic, page\_no, lower/upper).

## **ROW MODULE:**

### **encode\_row(Schema schema, List: row)**

Turns a row (Value per column) into raw bytes we can store on a page.

- Builds a **NULL bitmap** (bit  $i = 1 \Rightarrow$  column  $i$  is NULL).
- Appends payloads for **non-NULL** columns in schema order:
- Prepends a **TupleHeader** with totals (size, column count, nullmap length, flags).
- Returns a self-contained byte vector ready for Page::add\_tuple.

### **decode\_row(Schema schema, data, len)**

Reconstructs a row (Values) from the stored bytes.

- Reads and validates the **TupleHeader** (size & column count).
- Reads the **NULL bitmap**; for each NULL bit sets Value::null(type).
- For non-NULLs, parses payloads in schema order (same encodings as above).

TupleHeader{

total\_len; // full tuple bytes (hdr + bitmap + payload)

attr\_count; // number of columns

nullmap\_bytes; // bitmap length in bytes

flags; // bit0=dead (0 means live in Phase 0)

};

- Stored at the beginning of each tuple; used by both encode/decode.

## **HEAP MODULE:**

- **HeapTable::HeapTable(path, Schema schema)**

Opens (or creates) the heap file at path and remembers the table schema.

- **ensure\_first\_page()**

If the heap file is empty, appends page 0 and initializes it as a fresh slotted page (lower/upper pointers set).

- **Result insert(row)**  
Encodes the row (encode\_row), then tries to place it on the **last page**.
  - If there isn't enough space, appends a new page, initializes it, and inserts there.
  - Returns the tuple ID (**TID** = {page\_no, slot\_no}) on success.
- **Status remove(TID)**  
Loads the target page and marks the slot dead (tombstone).  
Then calls Page::compact() to reclaim space.  
Returns OK or NOT FOUND if the TID is invalid/unused.
- **scan()**  
**Full table scan:** for each page and each live slot, decodes the tuple and prints the tuples.
- **npages()**  
Returns current number of pages = filesize / PAGE\_SIZE.
- **bool read\_page(page\_no, out)**  
Seeks to page\_no \* PAGE\_SIZE and reads exactly one page into out.
- **bool write\_page(page\_no, buf)**  
Seeks to page\_no \* PAGE\_SIZE and writes one page from buf, then flushes.
- **append\_new\_page()**  
Extends the file by one zero-filled page and returns the new page number.

## CATALOG MODULE:

### What the Catalog does?

- Keeps a persistent mapping **table name → (relfilenode id, schema)**.
- Stores metadata in a human-readable file: data/catalog.txt.
- Creates/removes the per-table heap file under data/base/<relfilenode>.heap.
- **Catalog::Catalog(datadir)**  
Initializes a catalog rooted at datadir. Ensures datadir/ and datadir/base/ exist; sets catpath\_ = datadir/catalog.txt; starts next\_id at 1.
- **load()**  
Reads catalog.txt into memory.
  - Loads the TableEntry in tables.
- **create\_table(name, Schema s)**  
Allocates a new relfilenode (from next\_id++), records {name → (id, schema)}, and **touches** the heap file data/base/<id>.heap so it exists. Returns the new id.

- **bool drop\_table(name)**  
Removes name from the in-memory map and deletes its heap file data/base/<id>.heap. Returns true if the table existed.
- **heap\_path(relfilenode)**  
Builds the on-disk path for a table's heap file: datadir\_/base/<relfilenode>.heap.
- **ensure\_dir(p)**  
Creates a directory if missing (used for datadir/ and datadir/base/).

Data Members:

- datadir\_ – root directory (e.g., "data").
- catpath\_ – path to catalog.txt.
- next\_id – next relfilenode to assign.
- map<string, TableEntry> tables\_ – in-memory name→entry map.

## **SCHEMA:**

- **Column**
  - **Fields:** string name; Type type;
  - **Column(name, Type t)** — Constructor. Sets the column's name, type.
- **Schema**
  - **Fields:** Ordered list of columns that defines the row layout.
  - **size()** Returns the number of columns in the schema.

## **When you “start the DB”:**

1. **Construct the Catalog**
  - a. Catalog cat("data");
  - b. Constructor ensures directories exist:

data/

base/ (heap files live here)

catalog.txt (tiny metadata file)

- c. Sets `catpath_ = "data/catalog.txt"` and `next_id = 1`.

## 2. Load the catalog

- a. `cat.load();`
- b. Opens `data/catalog.txt` if it exists.
- c. Builds a Schema (columns with Type).

Records `tables_[name] = { relfilenode=id, schema }`.

So after startup we have:

- A memory map `name → {relfilenode, schema}`.
- A counter `next_id` for the next table id you'll allocate.

# Creating a table

## 1. Define a Schema

```
Schema s;  
s.cols.push(Column("id",  Type::INT);  
s.cols.push(Column("name", Type::TEXT);  
s.cols.push(Column("age",  Type::INT));  
s.cols.push(Column("active", Type));
```

## 2. Create the table in the catalog

- a. `rid = cat.create_table("people", s);`
- b. Inside `create_table`:
  - i. Assigns `relfilenode = next_id++`.
  - ii. Stores `tables_["people"] = { rid, schema }`.
  - iii. Touches the heap file so it exists: `data/base/1.heap` (empty file).
- c. `cat.save();` writes:

## 3. Open the heap file

- a. Look up the entry: `cat.lookup("people", entry)`.
- b. Construct the heap:

```
HeapTable heap(cat.heap_path(entry.relfilenode), entry.schema);
```

- c. HeapTable opens data/base/1.heap .  
(It does **not** allocate the first page yet)

## Inserting a row

Insert (1, "T", 26):

### Row in memory

```
row = {  
  Value::i v(1),    // id  
  Value::s v("T"),  // name  
  Value::i v(26),   // age  
};
```

### 1) HeapTable::insert

- heap.insert(row) does:
  - **Ensure first page exists**
    - ensure\_first\_page() checks file size.
      - If empty: append\_new\_page() writes 8192 zero bytes (page 0), then:
        - Creates a Page, p.init(0) sets:
          - header.magic
          - header.page\_no = 0
          - header.lower = sizeof(PageHeader) (start of ItemId array)
          - header.upper = PAGE\_SIZE (8192)
        - store\_to() → write initialized page 0 back to disk.
    - **Encode the row to bytes**
      - Calls encode\_row(schema, row):
        - Builds a NULL bitmap (1 bit per column). Appends payloads in schema order:
          - Id, name, age
        - Prepends a **TupleHeader** (packed):
          - total\_len = header + nullmap + payload

- attr\_count = 3
  - nullmap\_bytes = 1
  - flags = 0 (live)
  - Returns one contiguous byte vector representing the tuple.
- **Pick a page (last page) and add the tuple**
  - Reads page 0 into memory, constructs Page p from raw bytes.
  - Calls p.add\_tuple(tuple\_bytes):
    - Checks **free space**: upper - lower must be  $\geq$  tuple\_size + sizeof(ItemId).
    - **Places payload at the top**:
      - upper -= tuple\_size
      - copies tuple bytes to data\_[upper .. upper+tuple\_size)
    - **Appends an ItemId** at the end of the item array:
      - Computes current slot = nititems()
      - Writes ItemId{ offset=upper, length=tuple\_size, flags=ITEM\_USED }
      - lower += sizeof(ItemId)
    - Returns the **slot number** (e.g., 0 on the first insert).
  - HeapTable writes the modified page back to disk.
  - Returns the **TID**: { page\_no=0, slot\_no=0 }.

## Deleting a row

- heap.remove(TID):
  - Reads the page, p.mark\_dead(slot) sets ITEM\_DEAD in that slot's ItemId.
  - Calls p.compact():
    - Collects all **live** tuples.
    - Resets upper = PAGE\_SIZE.
    - Rewrites live payloads tightly from the top, updates their ItemIds (same slot numbers).
    - Clears dead slots (flags=0, length=0, offset=0).
  - Writes the page back.