


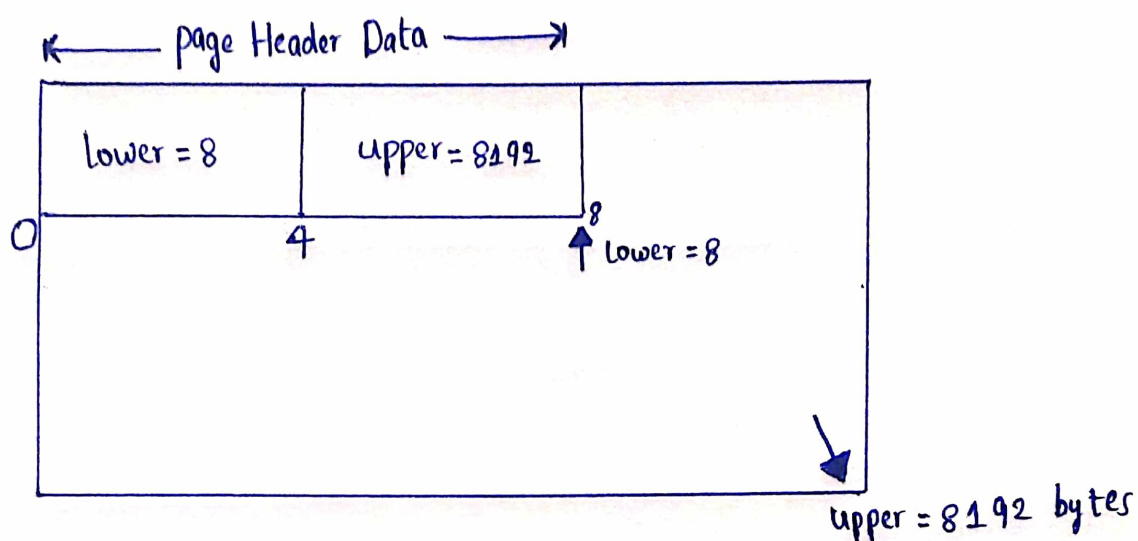
# **Storage Manager Design Doc - version 0**

## File/ Table Logical Layout



↓  
pub struct TableHeader {  
    pub page-count: u32 // Total pages in table  
}

## Initial page Data



Page = 8 KB

page Header = (lower - 4 bytes + upper - 4 bytes) = 8 bytes

lower pointer = 8<sup>th</sup> byte  
upper pointer = 8192<sup>th</sup> byte

## Table - Physical Layout

```
pub struct Table {  
    pub data: Vec<u8>, // Fixed-size buffer holds the raw bytes of a table.  
}
```

## Table - Logical Layout

```
pub const TABLE_HEADER_SIZE: u32 = 8192;  
  
pub struct TableHeader {  
    pub page_count: u32, // Total Number of Pages in a Table  
}  
  
pub struct Table {  
    pub table_header: TableHeader,  
    // Pages are laid out consecutively after the table header on disk  
}
```

## Page - Physical Layout

```
pub const PAGE_SIZE: usize = 8192;  
  
pub struct Page {  
    pub data: Vec<u8> // Fixed-size buffer holds the raw bytes of a page (PAGE_SIZE = 8KB)  
}
```

## Page Header

```
pub const PAGE_HEADER_SIZE: u32 = 8; // Page Header Size - 8 bytes (4 for lower, 4 for upper)  
  
pub struct PageHeader {  
    pub lower: u32, // Offset to start of free space - 4 bytes  
    pub upper: u32, // Offset to end of free space - 4 bytes  
}
```

## Item/Tuple Details

```
pub const ITEM_ID_SIZE: u32 = 8;  
  
pub struct ItemId {  
    pub offset: u32, // Offset of the item/tuple  
    pub length: u32, // Length of the item/tuple  
}
```

## Logical Page Layout

```
pub struct Page {  
    pub header: PageHeader,  
    pub item_id_data: Vec<ItemId>,  
    // Tuples and their metadata are organized within the page after the header  
}
```

- **Table** and **File** are used interchangeably in the document. Both Represent same.
- 

## Currently Implemented API's

1. Init Table
2. Init Page
3. Page Count
4. Create Page
5. Read Page
6. Write Page
7. Page Free Space
8. Page Add Data

## Ongoing API's

1. Create Table
  2. Add Tuple
- 

### 0. init\_table API

#### Description:

- Initializes the **Table Header** by writing the **first page** (8192 bytes) into the table file with 0's. The first 4 bytes represent the **Page Count** (0),

#### Function:

```
pub fn init_table(file: &mut File)
```

#### Input:

file: File pointer to update Table Header.

#### Output:

Table header (first page) initialized with page\_count = 0 in the first 4 bytes and remaining bytes set to zero.

#### Implementation:

1. Move the file cursor to the beginning of the file.
2. Allocate a buffer of 8192 bytes (**TABLE\_HEADER\_SIZE**) initialized to zero.
3. Write the entire 8192-byte buffer (including the page count) to disk, marking the creation of the first table page.

#### Test Case:

1. Created a new file to simulate a fresh table.
  2. Initialized the table header using **init\_table** API, setting the page count to 0.
  3. Verified that the file size obtained from **metadata().len()** equals **TABLE\_HEADER\_SIZE** (8192 bytes), confirming that the entire header page was written successfully.
- 

## 1. init\_page API

### Description:

- Initializes the **Page Header** with two offset values for **In Memory Page**:
  - **Lower Offset** (PAGE\_HEADER\_SIZE) → bytes 0..4
  - **Upper Offset** (PAGE\_SIZE) → bytes 4..8

### Function:

```
pub fn init_page(page:&mut Page)
```

### Input:

page: **In Memory Page** to set Header - Lower and Upper Offsets.

### Output:

Page header updated with lower and upper offsets.

### Implementation:

1. Write the lower offset (PAGE\_HEADER\_SIZE) into the first 4 bytes of the page header (0..4).
2. Write the upper offset (PAGE\_SIZE) into the next 4 bytes of the page header (4..8).

### Test Case:

1. Created a new in-memory page with zeroed data.
  2. Initialized the page header using init\_page API, setting the lower and upper offsets.
  3. Checked whether the first 4 bytes were **PAGE\_HEADER\_SIZE** and the next 4 bytes were **PAGE\_SIZE**.
- 

## 2. page\_count API

### Description:

To get total number of pages in a file

### Function:

```
pub fn page_count(file: &mut File)
```

### Input:

file: file to calculate number of pages.

### Output:

Total number of pages present in the file.

### Implementation:

1. Use the **read\_page()** function to read the first page (page ID 0) from the file into memory.

2. Extract the **first 4 bytes** from the in-memory page buffer — these bytes represent the page count stored in the table header.
3. Return the first 4 bytes as page count.

**Test Case:**

1. Create a temp table file.
  2. Initialize it using `init_table()`.
  3. Call `page_count()` to verify it correctly reads 0.
- 

### 3. `create_page` API

**Description:**

Create a page in disk for a file.

**Function:**

```
pub fn create_page(file: &mut File)
```

**Input:**

file: file to create to a file

**Output:**

1. Create a page at the end of the file.
2. Update the File Header with **Page Count**.

**Implementation:**

1. Initializes a new page **in memory** using **init\_page** API (update page header - lower and upper).
2. Reads the **current page count** from the file using the **page\_count** API.
3. Moves the file cursor to the end of the file.
4. Writes the initialized in-memory page to the file and **updates the file header** by incrementing the page count stored in the first 4 bytes.

**Test Case:**

1. Verified using File Size, Page Count and Page Headers before and after creating the page using file metadata.
- 

### 4. `read_page` API

**Description:**

Reads a page from a disk/file into memory.

**Function:**

```
pub fn read_page(file: &mut File, page: &mut Page, page_num: u32)
```

**Input:**

file: file to read from,  
page: memory page to fill,  
page\_num: page number to read

**Output:**

Populates the given memory page with data read from the file.

**Implementation:**

1. Calculates the **offset** as **(page\_num \* PAGE\_SIZE)** and moves the file cursor to the correct position.
2. Reads data from that offset position up to **offset + PAGE\_SIZE** and copies it into the page memory.

**Cases Handled:**

1. Checks the file size and returns an error if the requested page does not exist in the file.

**Test Case:**

- Verified **read\_page** API correctly reads one full page — file size equals PAGE\_SIZE and data matches the original written content.
- 

## 5.write\_page API

**Description:**

Write a page from memory to disk/file.

**Function:**

```
pub fn write_page(file: &mut File, page: &mut Page, page_num: u32)
```

**Input:**

file: file to write,  
page: memory page to copy from,  
page\_num: page number to write

**Output:**

Writes the contents of the given memory page to the file at the specified page offset.

**Implementation:**

1. Calculates the **offset** as **page\_num \* PAGE\_SIZE** and moves the file cursor to the correct position.
2. copy the contents of the given memory page from offset to offset + PAGE\_SIZE positions to the file.

**Test Case:**

- Verified **write\_page** API correctly writes one full page — by using write\_page wrote some content to the page and read back the same page to confirm that the written data matches exactly, and the file size is at least PAGE\_SIZE.
- 

## 6. page\_free\_space API

**Description:**

To calculate the total amount of free space left in the page.

**Function:**

```
pub fn page_free_space(page: &Page)
```

**Input:**

page: page to calculate the free space.

**Output:**

Total amount of freespace left in the page.

**Implementation:**

1. Read the lower pointer from the first 4 bytes of the page.
2. Read the upper pointer from the next 4 bytes of the page.
3. Calculate free space = upper - lower.
4. Return the free space.

**Test Case:**

Verified that **page\_free\_space()** correctly calculates free space (upper - lower) for a newly created data page initialized using `init_table()` and `create_page()`.

Confirmed that the page header offsets (lower = `PAGE_HEADER_SIZE`, upper = `PAGE_SIZE`) and total page count are consistent and accurate.

---

## 7. page\_add\_data API

**Description:**

Adds raw data to the file.

**Function:**

```
pub fn page_add_data(file: &mut File, data: &[u8])
```

**Input:**

file: The file to which data should be added.

data: The raw bytes to insert into the page.

**Output:**

Data inserted in the file.

**Implementation:**

1. Get the **total number of pages** in the file using [page\\_count](#) API.
  2. Read the **last page** into memory using [read\\_page](#) API.
  3. Check **free space** in the page using [page\\_free\\_space](#) API.
  4. If the last page has enough free space to store the data and its ItemId (i.e., if `free_space >= data.size() + ITEM_ID_SIZE`):
    - a. Calculate the **insertion offset** from the upper pointer.  
`start = upper - data.len()`
    - b. Copy the data bytes into the page buffer starting at this offset.
    - c. Update the **upper pointer** in the page header to the new start of free space.
    - d. Write the **ItemId entry** (offset and length of the data) at the position indicated by the lower pointer.
    - e. Update the **lower pointer** in the page header to account for the newly added ItemId (`lower += ITEM_ID_SIZE`).
    - f. Write the updated page back to disk using [write\\_page](#) API.
  5. If the last page does not have enough free space:
    - a. [TODO]
-



- [Code - Github \(https://github.com/hemanth-sunkireddy/Storage-Manager\)](https://github.com/hemanth-sunkireddy/Storage-Manager)
- [Code Documentation \(https://hemanth-sunkireddy.github.io/Storage-Manager/storage\\_manager/all.html\)](https://hemanth-sunkireddy.github.io/Storage-Manager/storage_manager/all.html)
- **Reference 1:** API Formats – [Storage Manager Course Assignment Link \(http://www.cs.iit.edu/~glavic/cs525/2023-spring/project/assignment-1/\)](http://www.cs.iit.edu/~glavic/cs525/2023-spring/project/assignment-1/)
- **Reference 2:** [Postgres Internals – Page Layouts & Data \(https://www.postgresql.org/docs/current/storage-page-layout.html\)](https://www.postgresql.org/docs/current/storage-page-layout.html)